



Recap

2

- ▶ HDFS filesystem
 - Block
 - Read and Write operations
 - CLI commands
 - Network topology and rack awareness
- ▶ MapReduce

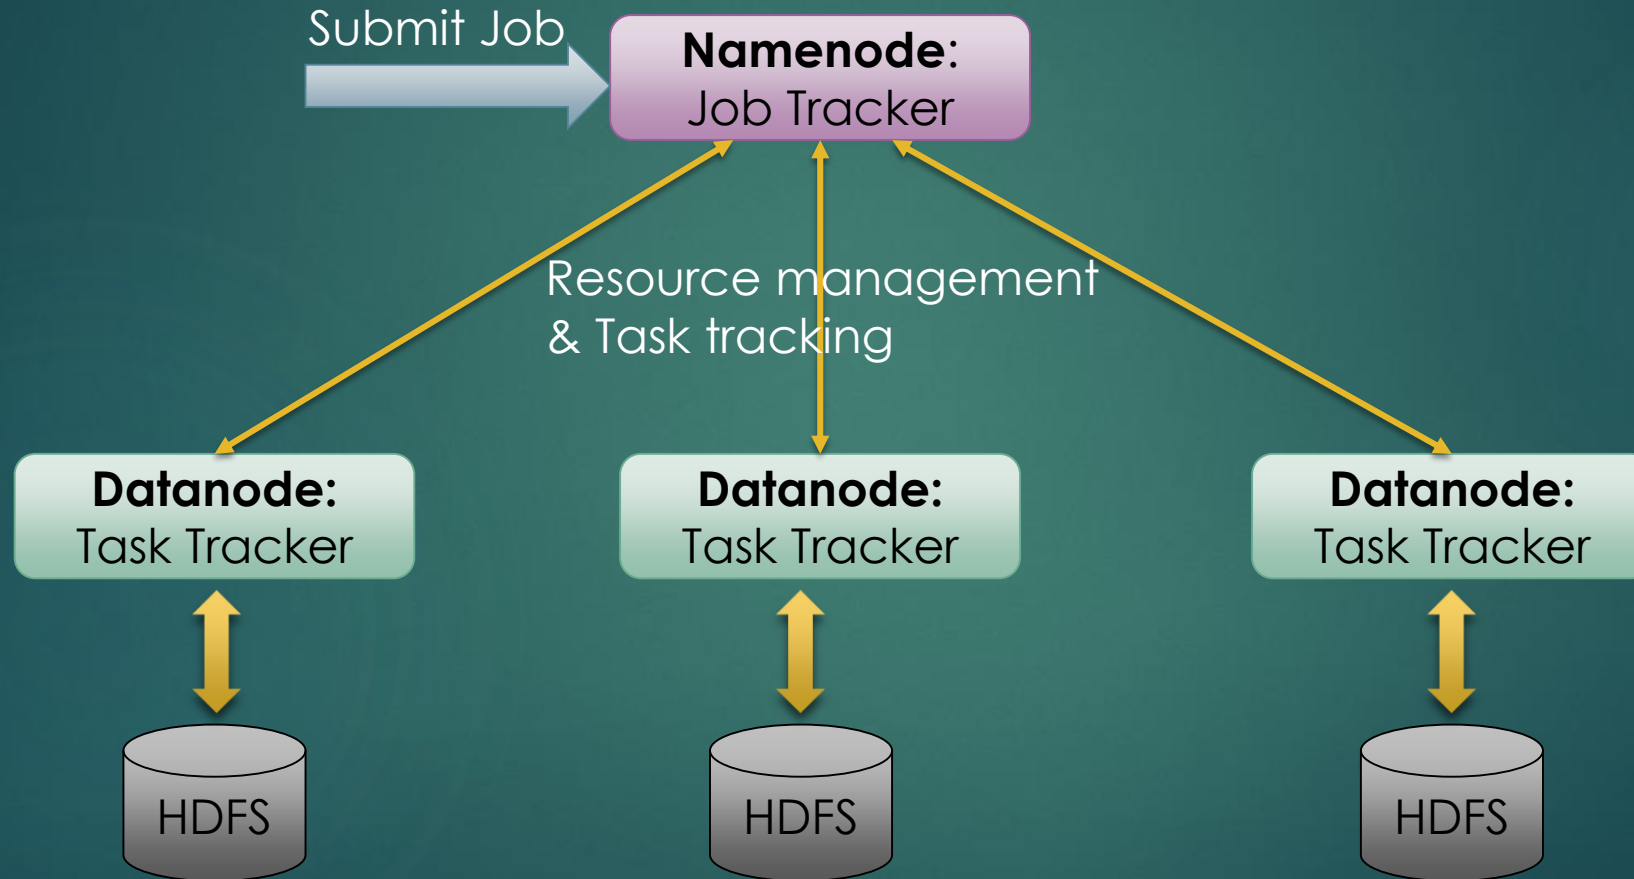
Agenda for today

3

- ▶ MapReduce detailed discussion
- ▶ Running your first MapReduce program
- ▶ Hadoop 1 vs 2: YARN

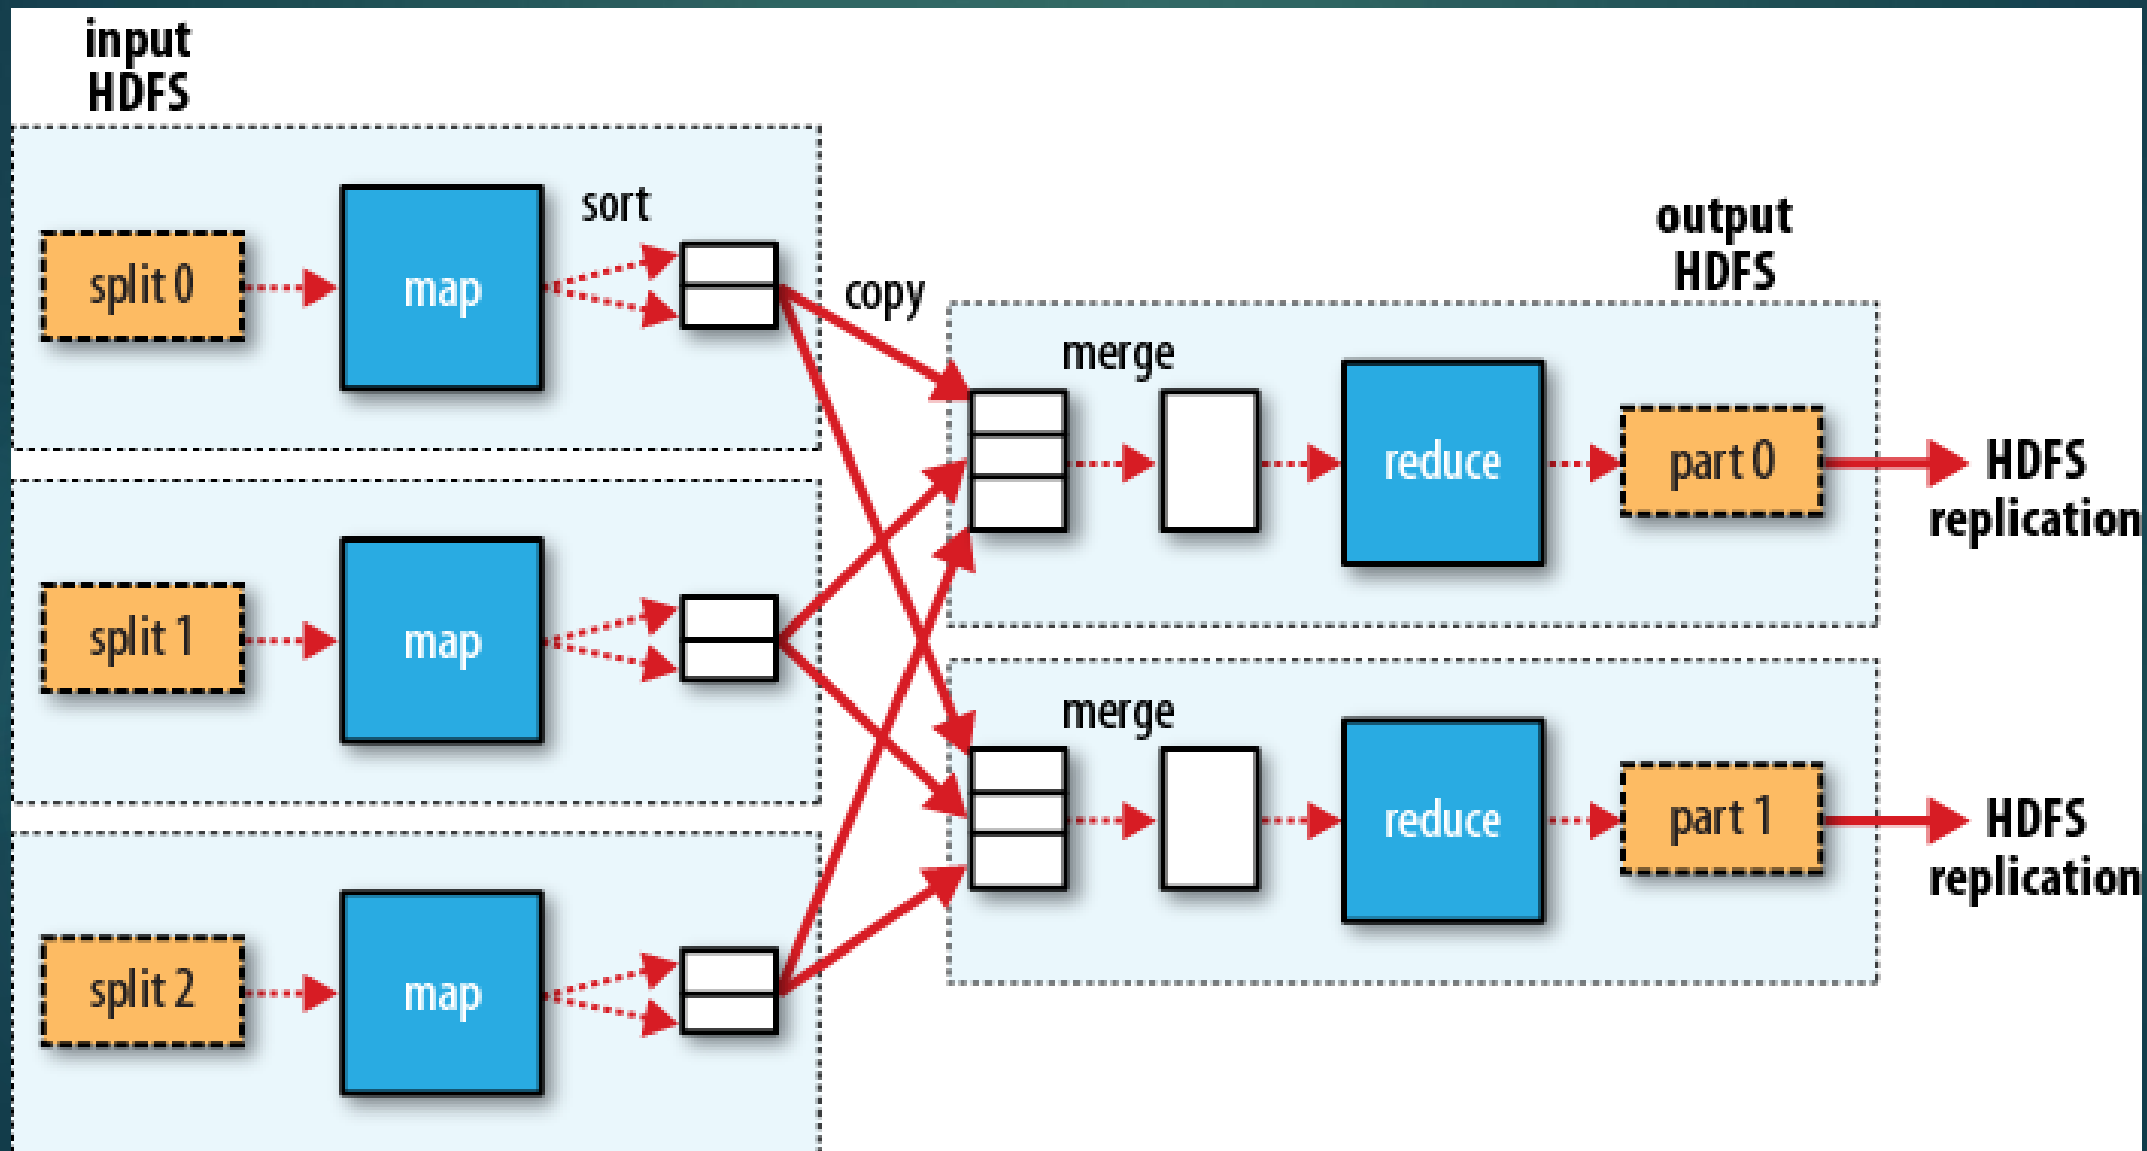
Job management

4



MapReduce Stages

5



Mapreduce: working example

6

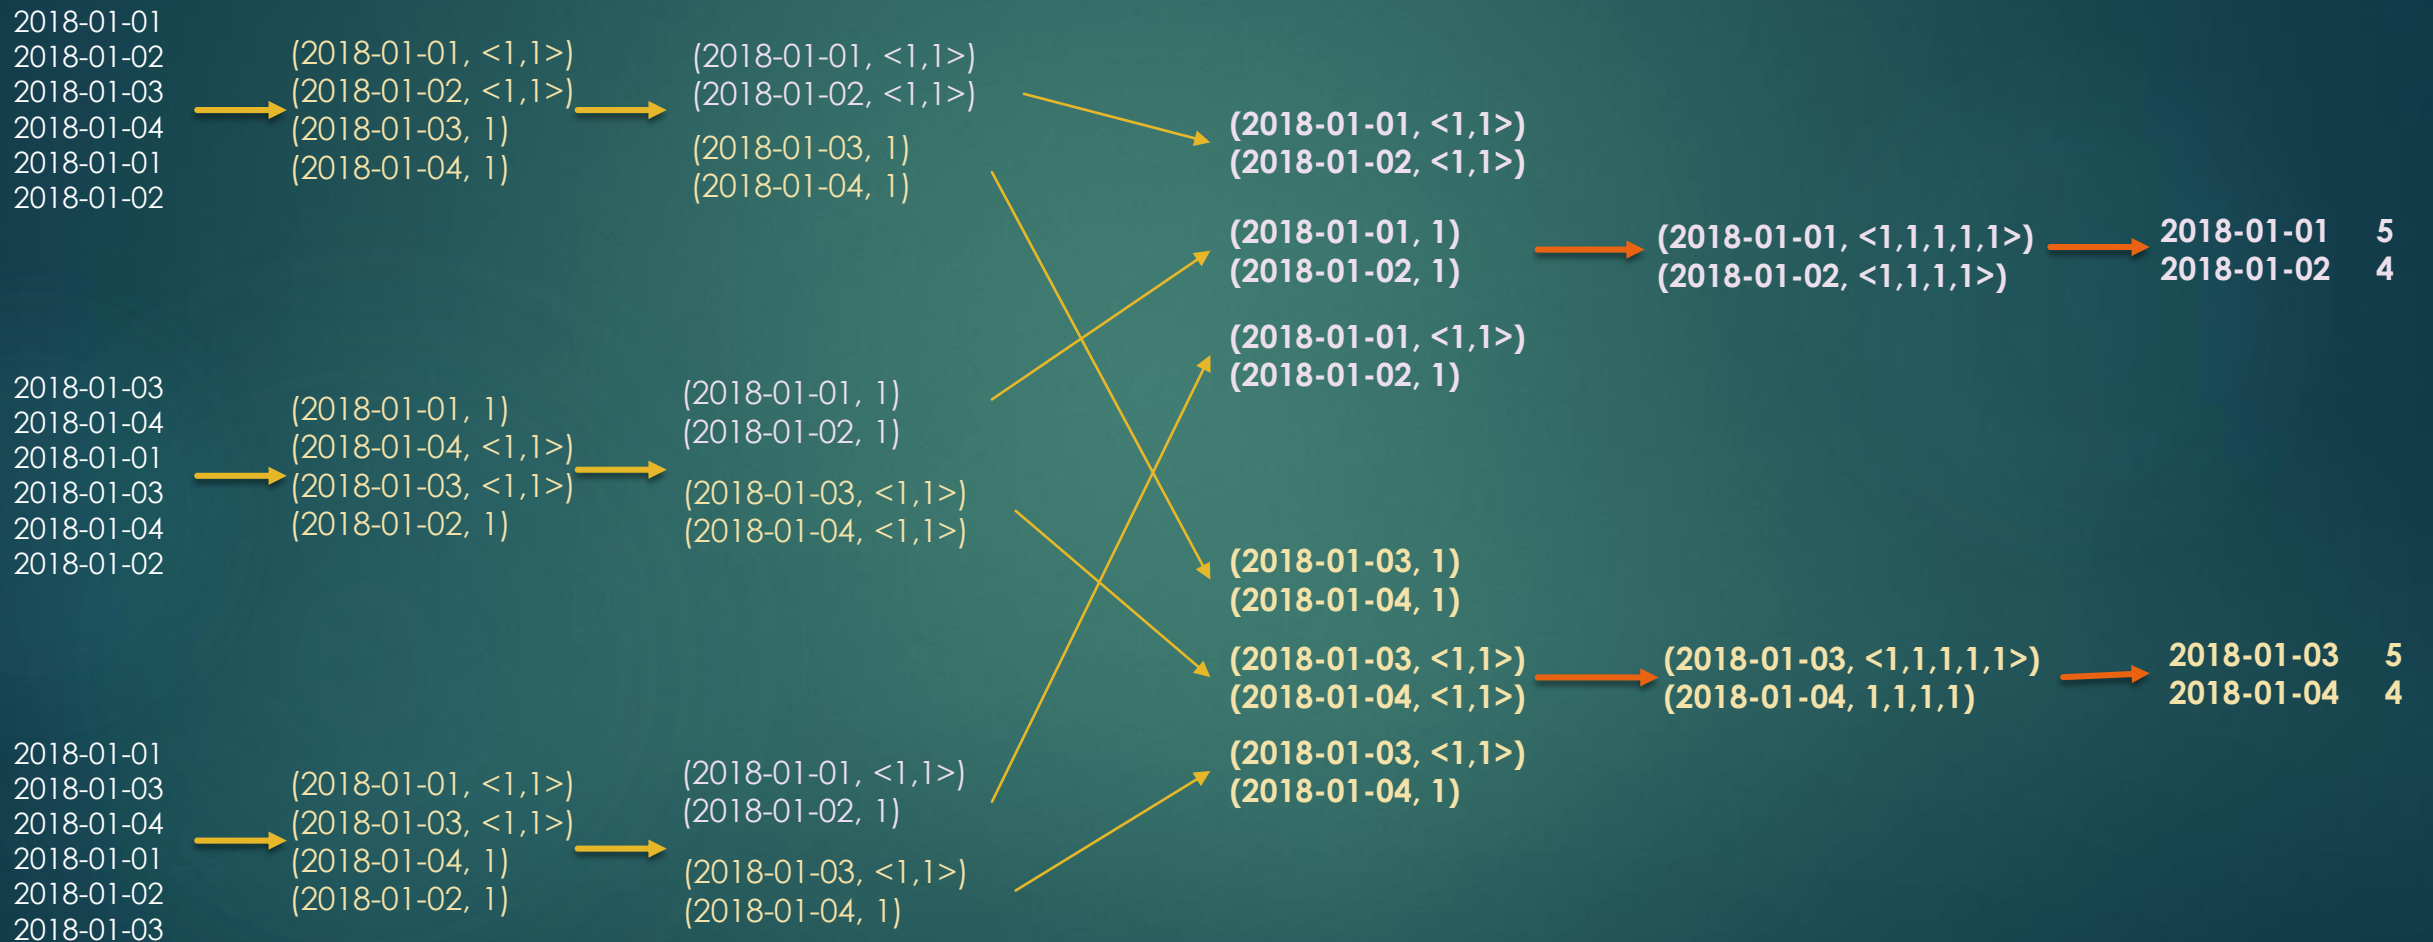
Map

Sort & Partition

Shuffle


Merge

Reduce



Task to node mapping


- Notion of data locality



Input Splits

8

- ▶ Blocks are of fixed size
- ▶ Good chances of records being split between two block



MapReduce: Mapper code

```
public class WebHitCounterMapper extends
Mapper<Input Key, Input Value, Output Key, Output Value>
{

    public void map(Input Key, Input Value, Context context)
throws IOException, InterruptedException {

        <MAP Logic goes here>

        context.write(Output Key, Output Value)
    }
}
```

MapReduce: Reducer code

```
public class WebHitCounterReducer extends
Reducer<Input Key, Input Value, Output Key, Output Value>
{
    public void reduce(Input Key, Iterable<Value Data type>
values, Context context) throws IOException,
InterruptedException {

    <REDUCE logic goes here>

    context.write(Output Key, Output Value);
}
}
```

MapReduce: Driver Code

11

```
public class WebHitCounterMain {  
    public static void main(String[] args) throws Exception {  
  
        Configuration conf = new Configuration();  
        Job job = Job.getInstance(conf, "Daily Web Hit Counter");  
  
    }  
}
```

MapReduce: Driver Code

12

```
public class WebHitCounterMain {  
    public static void main(String[] args) throws Exception {  
  
        Configuration conf = new Configuration();  
        Job job = Job.getInstance(conf, "Daily Web Hit Counter");  
  
        job.setJarByClass(main.WebHitCounterMain.class);  
        job.setMapperClass mapper.WebHitCounterMapper.class);  
        job.setReducerClass(reducer.WebHitCounterReducer.class);  
  
    }  
}
```

MapReduce: Driver Code

13

```
public class WebHitCounterMain {  
    public static void main(String[] args) throws Exception {  
  
        Configuration conf = new Configuration();  
        Job job = Job.getInstance(conf, "Daily Web Hit Counter");  
  
        job.setJarByClass(main.WebHitCounterMain.class);  
        job.setMapperClass mapper.WebHitCounterMapper.class);  
        job.setReducerClass(reducer.WebHitCounterReducer.class);  
  
        job.setOutputKeyClass(Text.class);  
        job.setOutputValueClass(IntWritable.class);  
  
    }  
}
```

MapReduce: Driver Code

14

```
public class WebHitCounterMain {
    public static void main(String[] args) throws Exception {

        Configuration conf = new Configuration();
        Job job = Job.getInstance(conf, "Daily Web Hit Counter");

        job.setJarByClass(main.WebHitCounterMain.class);
        job.setMapperClass mapper.WebHitCounterMapper.class);
        job.setReducerClass(reducer.WebHitCounterReducer.class);

        job.setOutputKeyClass(Text.class);
        job.setOutputValueClass(IntWritable.class);

        FileInputFormat.addInputPath(job, new Path(args[0]));
        FileOutputFormat.setOutputPath(job, new Path(args[1]));

    }
}
```


MapReduce: Driver Code

15

```
public class WebHitCounterMain {  
    public static void main(String[] args) throws Exception {  
  
        Configuration conf = new Configuration();  
        Job job = Job.getInstance(conf, "Daily Web Hit Counter");  
  
        job.setJarByClass(main.WebHitCounterMain.class);  
        job.setMapperClass(mapper.WebHitCounterMapper.class);  
        job.setReducerClass(reducer.WebHitCounterReducer.class);  
  
        job.setOutputKeyClass(Text.class);  
        job.setOutputValueClass(IntWritable.class);  
  
        FileInputFormat.addInputPath(job, new Path(args[0]));  
        FileOutputFormat.setOutputPath(job, new Path(args[1]));  
  
        System.exit(job.waitForCompletion(true) ? 0 : 1);  
    }  
}
```

Programming Exercise

16

Challenges with Hadoop 1

17

- ▶ Applications were limited to MapReduce implementations only
- ▶ Namenode machine crash or maintenance activity
- ▶ Namespace scaling
- ▶ Backup and Recovery
- ▶ Batch oriented architecture
- ▶ Support for various file formats
- ▶ Dual responsibilities of Job tracker

Image Ref: <https://www.greycampus.com/blog/big-data/top-differences-between-hadoop-1-0-and-hadoop-2-2>

Hadoop 2

18

- ▶ Support for other data processing engines
- ▶ High Availability
- ▶ HDFS Federation
- ▶ HDFS Snapshot
- ▶ Introduced Streaming and Interactive analysis
- ▶ Support for various file formats
- ▶ Yarn

► Yet Another Resource Negotiator

MapReduce 1	YARN
Job Tracker	Resource Manager, Application Master and Timeline server
Task Tracker	Node Manager
Slot	Containers

YARN model

20

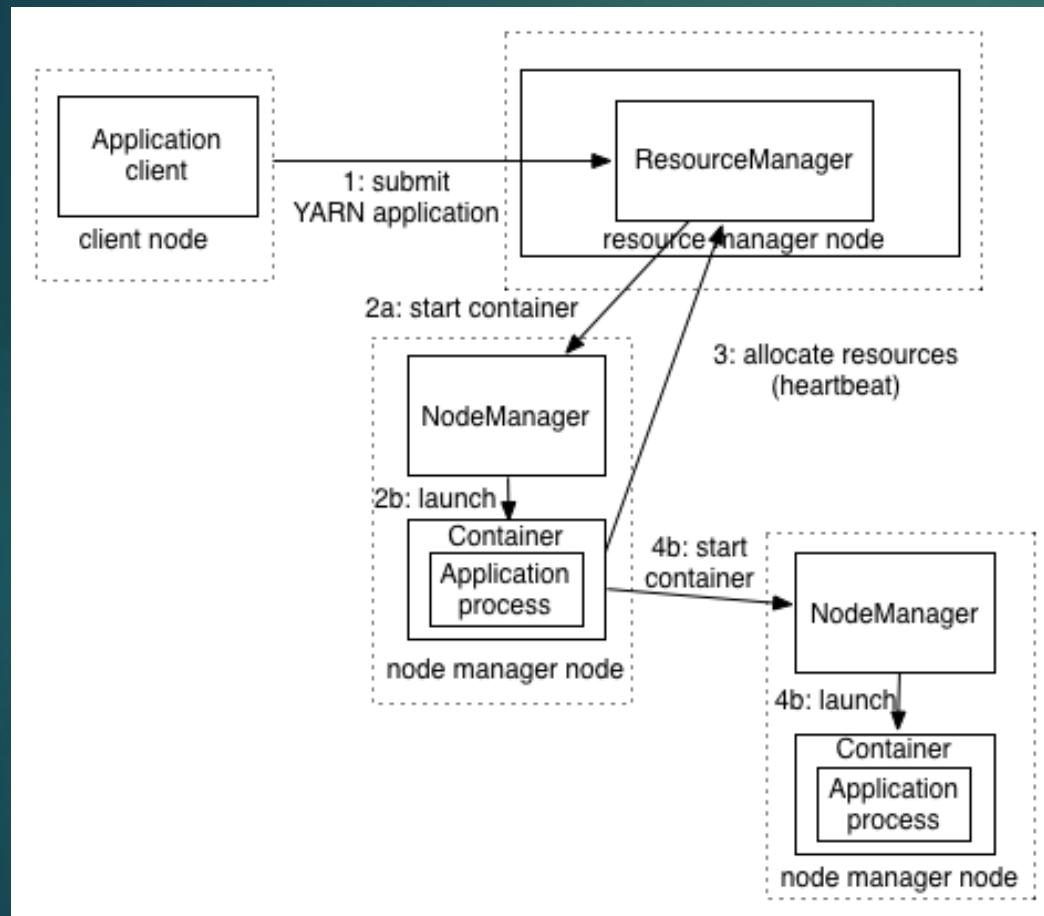


Image Ref: Hadoop definitive guide 4th edition

Pros of YARN

21

- ▶ Scalability
- ▶ Availability
- ▶ Utilization
- ▶ Multitenancy

Hadoop Installation

22

Standalone

Everything in one JVM. No HDFS installation

Pseudo-distributed

Mimic a distributed cluster on single physical machine

Distributed

Fully distributed cluster with multiple physical machines

Reference

23

- ▶ Hadoop standalone vs pseudo-distributed

<https://stackoverflow.com/questions/23435333/what-is-the-difference-between-single-node-pseudo-distributed-mode-in-hadoop>

- ▶ Hadoop installation differences

<https://medium.com/@nidhinmahesh/getting-started-hadoop-mapreduce-hdfs-and-yarn-configuration-and-sample-program-febb1415f945>

- ▶ Download Ubuntu

<https://www.ubuntu.com/download/desktop>

- ▶ Hadoop installation step by step guide

http://www.bogotobogo.com/Hadoop/BigData_hadoop_Install_on_ubuntu_16_04_single_node_cluster.php