



# Recap

2

- ▶ Hadoop installation
- ▶ HDFS Java API

# Agenda for today

3

- ▶ MapReduce counters
- ▶ Performance tuning in MapReduce jobs
- ▶ MapReduce job chaining
- ▶ Pig
- ▶ Java programming & SQL for beginners

# Performance tuning

4

- ▶ Cluster configuration
- ▶ Use compression technique
- ▶ Tuning # mappers and reducers
- ▶ Use combiner
- ▶ Appropriate data type
- ▶ Reuse objects
- ▶ Profiling

<https://blog.cloudera.com/blog/2009/12/7-tips-for-improving-mapreduce-performance/>



# MapReduce Job chaining

5

- ▶ Two separate jobs
- ▶ Multiple mappers/reducers within same job

# MapReduce Job chaining

6

- ▶ Two separate jobs

1. Configure first job object and run it.
2. Configure second job object and run it

# MapReduce Job chaining

7

- ▶ Multiple mappers/reducers within same job
  - ▶ ChainMapper API: to add multiple mappers
  - ▶ ChainReducer API: to add multiple reducers



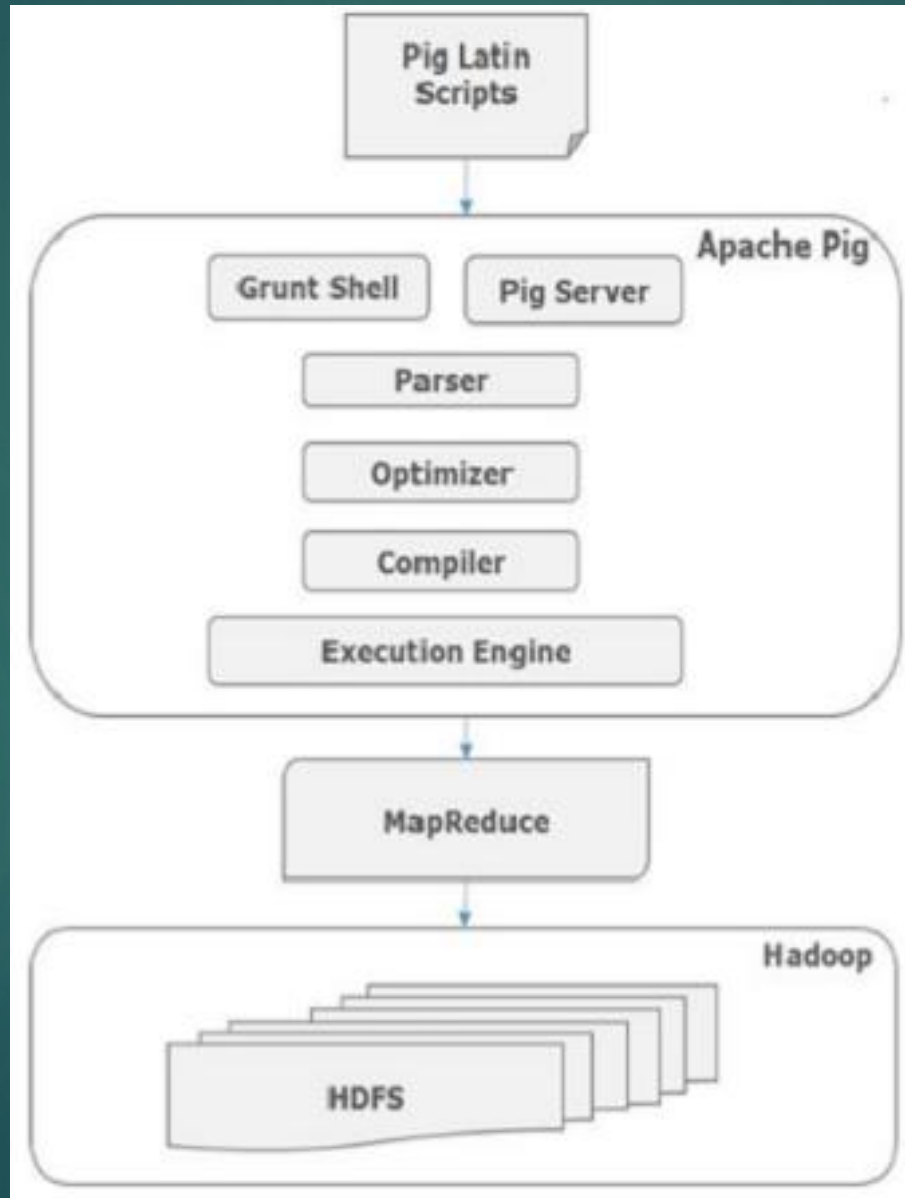


# Introduction

- ▶ High Level Scripting Language developed by Yahoo originally
- ▶ Transforms SQL like language called Pig Latin into Java code
- ▶ Follows lazy evaluation
- ▶ Supports UDF written in multiple languages

# Architecture

10



# Execution

11

## ▶ Accessing approaches:

1. Batch mode: submit a script directly
2. Interactive mode: Grunt, the pig shell
3. PigServer for Java program

## ▶ Execution mode:

1. Local mode: `pig -x local`
2. Mapreduce mode (default): `pig -x mapreduce`

# Data types

12

- ▶ Scalar Types: Int, long, float, double, boolean, null, chararray, bytearray
- ▶ Complex Types: fields, tuples, bags, relations

# Operator: LOAD

13

- ▶ To load data from storage system

```
lines=LOAD 'myfile' AS (line: chararray);
```

```
books = LOAD '/data/pig/books.csv' as (line: chararray)
```



# Operator: DUMP

14

- ▶ Print the data on console

DUMP RelationName;

DUMP sample\_books;

# Various loaders

15

► Supports various loader formats

1. TextLoader
2. PigStorage
3. JsonLoader & JsonStorage
4. BinStorage
5. HBaseStorage
6. OrcStorage
7. MongoStorage

# Operator: LOAD cont...

16

- ▶ Load data without schema

```
relXYZ = LOAD 'yourfile.csv' USING PigStorage(',');
```

```
books = LOAD '/data/pig/books.csv' USING PigStorage(',');
```

- ▶ Load data with schema

```
relXYZ = LOAD 'yourfile.csv' USING PigStorage(',') as  
(col1:datatype, col2:datatype,...);
```

```
books = LOAD '/data/pig/books.csv' USING PigStorage(',') as (id:int,  
author:chararray, name:chararray, year:int);
```

# Operator: LIMIT

17

- ▶ Take sample records

New\_Rel = LIMIT RelationName <Sample Count>;

Sample\_books = LIMIT books 5;

# Operator: FOREACH

18

- ▶ Select specific columns

New\_Rel = FOREACH RelationName GENERATE  
col1, col2, col3....;

book\_no\_author = FOREACH books GENERATE id, name, year;



# Operator: JOIN

19

- ▶ Joins two relations/datasets

```
join_data = JOIN relation1 BY (column1), relation2  
BY (column1);
```

```
book_review = JOIN books BY (id), reviews BY (id)
```

# Operator: SORT

20

- ▶ Sort a relation based on key

New\_rel = ORDER RelationName BY ColumnName  
asc;

books\_sorted\_by\_year = ORDER books BY year asc;

# Operator: FILTER

21

- ▶ Filter the dataset

New\_rel = FILTER RelationName BY (Condition);

books\_before\_2000 = FILTER books BY (year < 2000)

# Operator: DISTINCT

22

- ▶ Remove duplicates

New\_rel = DISTINCT RelationName;

bedupe = DISTINCT books\_before\_2000;

# Aggregate

23

- ▶ Aggregate based on a key

GroupRel = GROUP RelName BY columnName;

AggRel = FOREACH GroupRel GENERATE group ,  
AVG(columnName)

```
group_review = group book_review by books::id;
```

```
avg_rating = foreach group_review generate group as id,  
AVG($1.reviews::rating)
```



# Operator: STORE

24

- ▶ Store the output

```
STORE relationName INTO 'output_directory' USING  
PigStorage(',');
```

```
STORE dedupe INTO '/data/pig/dedupe' USING PigStorage(',');
```

# PigServer API

25

```
import java.io.IOException;
import org.apache.pig.PigServer;

public class idlocal{

    public static void main(String[] args) {
        try {
            PigServer pigServer = new PigServer("local");
            runIdQuery(pigServer, "passwd");
        }
        catch(Exception e) {}
    }

    public static void runIdQuery(PigServer pigServer, String inputFile) throws IOException {
        pigServer.registerQuery("A = load " + inputFile + " using PigStorage(':');");
        pigServer.registerQuery("B = foreach A generate $0 as id;");
        pigServer.store("B", "id.out");
    }
}
```

# UDF

26

- ▶ Prepare a Jar file
- ▶ Register the Jar
- ▶ Define alias
- ▶ Use it

[https://www.tutorialspoint.com/apache\\_pig/apache\\_pig\\_user\\_defined\\_functions.htm](https://www.tutorialspoint.com/apache_pig/apache_pig_user_defined_functions.htm)

# Further reading

27

- ▶ Map Reduce job chaining:

<https://mapr.com/blog/how-to-launching-mapreduce-jobs/>

- ▶ Pig inbuilt functions

<https://pig.apache.org/docs/latest/func.html>